MANGAL ANALYTICS AND® RESEARCH CONSULTING

Delivering Excellence, Partnering Success.

MARC Insights
The AI Capex Paradox
A Case Study

Is the 2026 hyperscaler buildout durable structural growth - or a classic capital cycle heading for overbuild?

August 2026

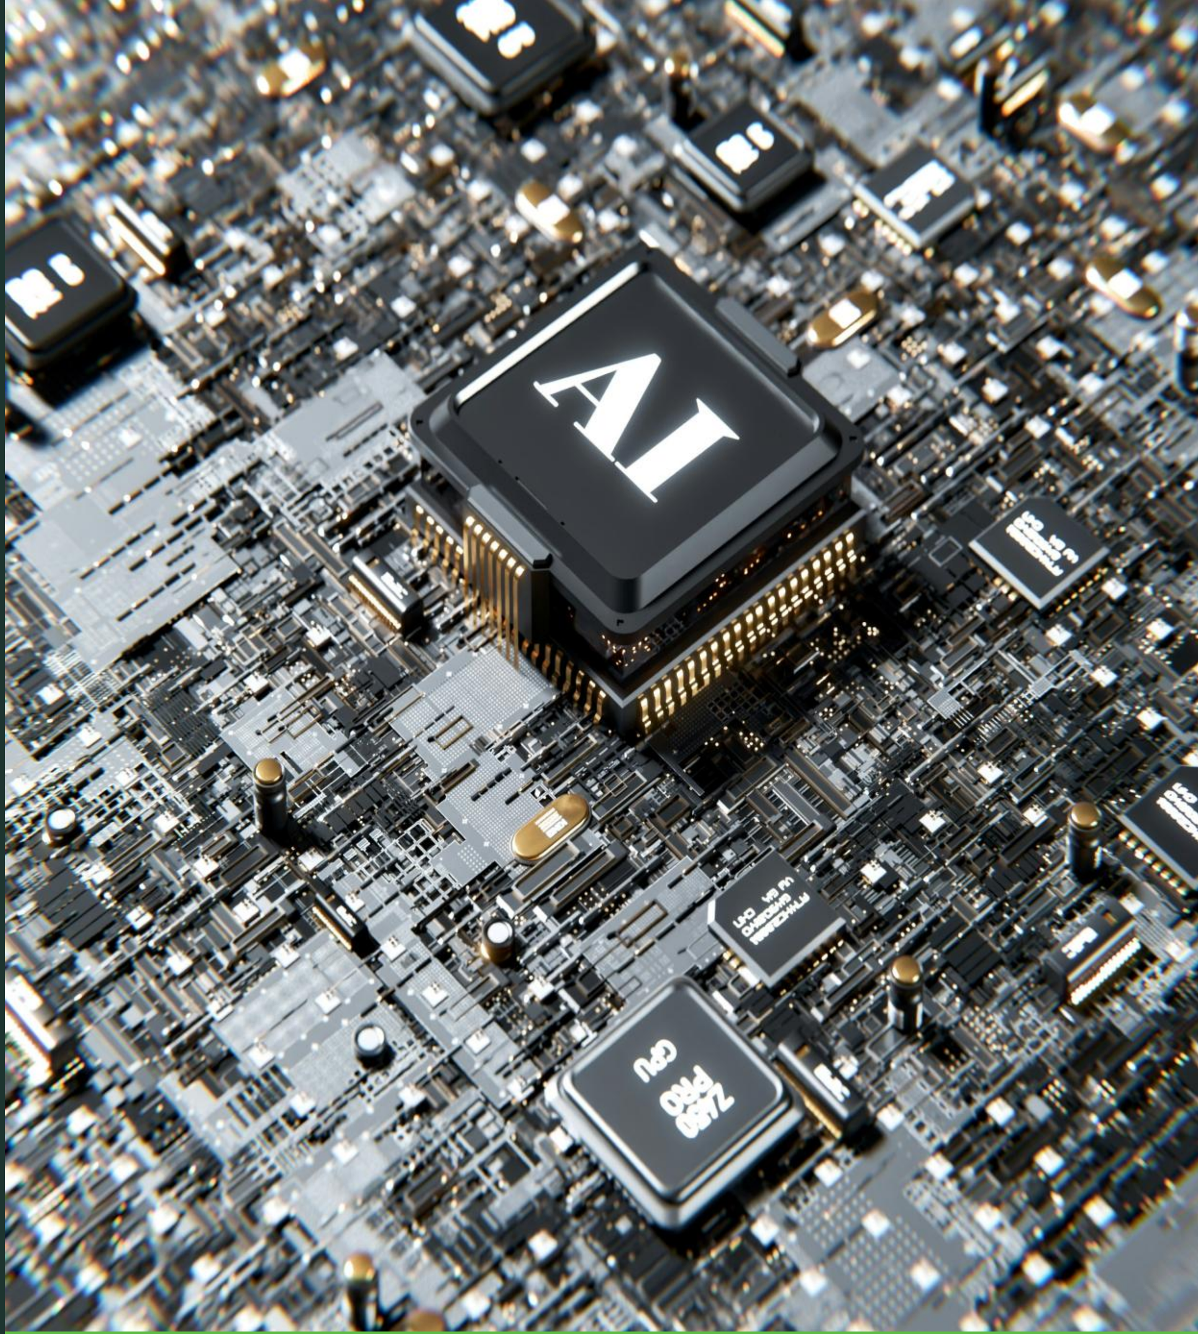

Table of Contents

SECTION 01

The Bill Arrives

Mid-2026. Two signals surfaced within weeks of each other - one at the retail shelf, one on the trading screen. Neither was about artificial intelligence directly. Both were the first visible invoices for it.

The Story Begins at a Checkout Counter

On 26 June 2026 Apple raised list prices across almost its entire hardware range. The stated reason was not tariffs, design or demand - it was memory. "The rapid expansion of AI data centres has created an extraordinary surge in demand for memory and storage."

Announced price increases - Apple hardware, June 2026

Product	Before	After	Change
Apple TV	$129	**$199**	+54%
iPad Pro	$999	**$1,199**	+20%
MacBook Air	$1,099	**$1,299**	+18%
Mac Studio (M3 Ultra)	$3,999	**$5,299**	+33%
HomePod mini	$99	**$129**	+30%
iPad Air	$599	**$749**	+25%

It was not only Apple

- Lenovo, Dell and HP all repriced through mid-2026, citing the same component pressure.
- DRAM chip prices roughly doubled in Q1 2026 alone.
- IDC expects the average selling price of a PC to rise 18.3% across 2026.

"Demand surges in one market are spilling over into cost surges in another."

Prof. Srikanth Jagabathula, NYU Stern - on the "hidden AI tax"

Why it matters Households began paying for the AI buildout before any of them bought an AI product. That is the first evidence that this capex cycle has consequences outside the balance sheets financing it.

Sources Al Jazeera, AppleInsider, MacRumors, TrendForce, The Register, IDC, Fortune

Weeks Later, the Market Flinched

The same memory economics that lifted chip valuations through 2025 turned on the sector in mid-2026. Two waves of selling - late June and late July - repriced the trade that had funded the buildout.

~$1.4 Tn
Value erased across AI chip names
Two selling waves, June-July 2026

-10%
Philadelphia Semiconductor Index
Single-session fall, June 2026

-$279 Bn
Nvidia market value
In the June drawdown alone

~-30%
Micron, SK Hynix, SanDisk
Peak-to-trough by late July 2026

What actually moved

- Memory names fell hardest, reversing the trade that had priced HBM scarcity as permanent.
- SoftBank, a proxy for leveraged AI exposure, sold off with the group.
- Several technology IPOs were pushed back as issuers waited for pricing conditions to settle.

What it told investors

- Markets stopped treating every dollar of AI capex as automatically value-accretive.
- The sell-off hit the supply chain harder than the hyperscalers commissioning the spend.
- It was a stress test, not a collapse: the sector recovered part of the loss within weeks.

Why it matters Two invoices arrived in the same quarter - one to consumers, one to shareholders. Neither settles the argument, but together they mark the point at which the AI buildout stopped being cost-free to everyone outside it.

Sources *Reuters / NST, CNBC, CryptoBriefing., 24/7 Wall St., CNBC, Reuters/Investing.com*

One Set of Facts, Two Opposing Conclusions

Everything that follows - the spending, the shortages, the revenue shortfall, the historical parallel - is read one way by each side of this argument. The purpose of this study is to set both cases out honestly and identify what would settle them.

Reading A - Durable structural growth

Demand is real and compounding

Global data volume is on track from 2 zettabytes in 2010 to roughly 221 zettabytes in 2026. Compute demand follows data.

The spenders can afford it

Unlike 2000, the buildout is funded largely from the operating cash flow of profitable businesses, not speculative debt.

Monetisation is early, not absent

Frontier model revenue has scaled from near zero to roughly $75 billion in run-rate in under three years.

Under-building is the greater risk

In a winner-take-most compute market, arriving late with capacity is more costly than arriving early with too much.

Reading B - A classic capital cycle

The hurdle is not close to being met

JPMorgan puts the revenue needed for a baseline 10% return at $650 billion a year. The industry earns roughly $75 billion.

The buyer is not converting

McKinsey, BCG and MIT independently find most enterprise AI deployments are not delivering the returns projected.

The spend leaves no margin for error

PIMCO analysis shows Big Tech directing about 94% of operating cash flow into AI capex.

The pattern is well documented

Capital inflow, simultaneous capacity expansion, yield collapse, consolidation - the fibre build of 1996-2002 followed it exactly.

MARC view

MARC's position: both readings are internally consistent today. What separates them is not conviction but time - specifically, whether enterprise monetisation closes the gap before the depreciation on this capex starts to bite.

***Sources** Exploding Topics, Source: 24/7 Wall St. J.P. Morgan AM BCG (Slide 20) PIMCO*

SECTION 02

The Build

To understand why a memory chip repriced a laptop, follow the money backwards. In six years the four largest cloud operators moved from spending $90 billion a year on infrastructure to committing $725 billion - an eightfold increase with no direct precedent in corporate history.

Eight Times More Capital in Six Years

Combined capital expenditure by Amazon, Alphabet, Meta and Microsoft. The curve did not bend gradually - it inflected in 2024 and then doubled twice.

Combined hyperscaler capital expenditure, 2020-2026

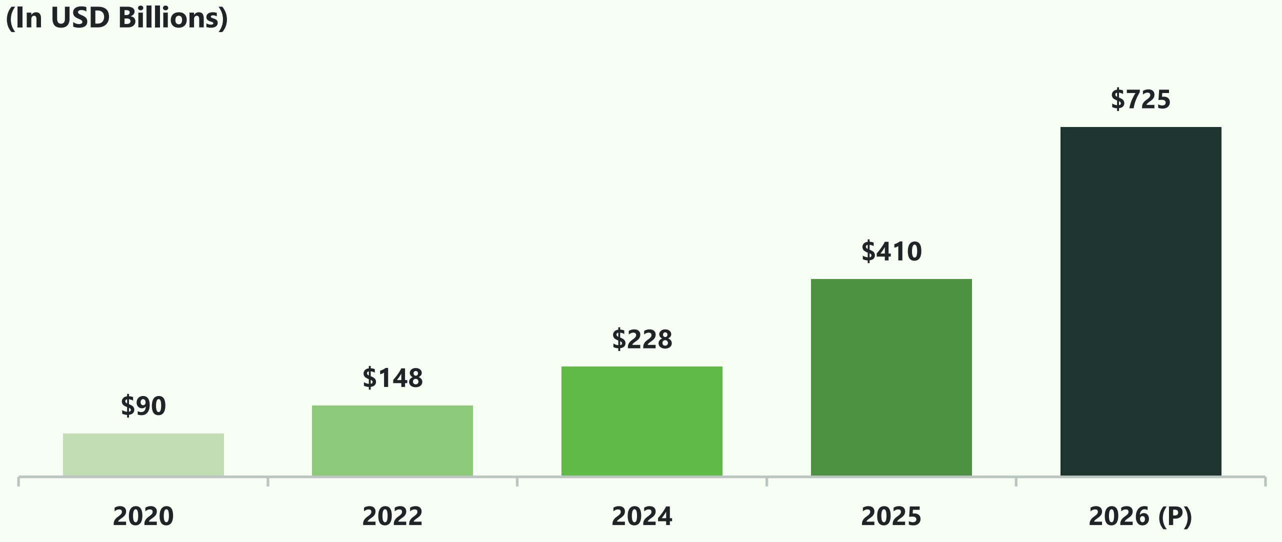

What the curve is actually saying

- The step change came in 2024-25, when GPU fleets and mega-facilities began scaling in parallel rather than in sequence.
- 2026 guidance alone (+77% year on year) is larger than the entire 2020-2023 period combined.
- No player can opt out: capacity is being added simultaneously by four competitors, each afraid of arriving late.
- The debate is live on the sell side. Jefferies calls the bear case "garbage"; SLC Management warns the spend could turn "a capital-light money machine into a capital-intensive incinerator."

Why it matters Simultaneous capacity expansion by well-funded competitors is precisely the condition under which capital cycles form. It is also precisely how genuine platform shifts get built. The curve alone cannot tell you which one this is.

Sources *Tom's Hardware,, MARC analysis*

Nearly All the Cash Flow Is Going One Way

Capex is affordable in absolute terms. The question is what share of internally generated cash it absorbs - and on that measure there is very little headroom left.

Share of Big Tech operating cash flow directed to AI capex

94%

of operating cash flow

is being redirected into AI infrastructure, per PIMCO analysis

Remaining 6% covers dividends, buybacks, M&A and every other use of capital

The three levers if demand disappoints

Cut capex
Immediate earnings relief; cedes capacity position to rivals

Cut returns
Protects the buildout; directly penalised by equity holders

Raise capital
Preserves both; reintroduces 2000-style leverage

Why the ratio matters more than the total

- At 94%, the buildout is self-funding - which is the strongest argument against a 2000-style credit event.
- It is also the weakest point of flexibility: there is no cushion left if AI revenue disappoints.
- Any shortfall must be absorbed by cutting capex, cutting shareholder returns, or raising external capital - all three of which markets read as a signal.
- Depreciation on assets already built begins hitting reported earnings well before the revenue case is settled.

Why it matters A buildout financed from cash flow cannot go bankrupt. It can, however, quietly destroy a great deal of shareholder value - which is the more likely failure mode this time.

Sources *PIMCO*

What $10-25 Billion Actually Buys

Aggregate capex figures are hard to interpret. Broken down to a single modern facility, the arithmetic becomes concrete - and shows how much of the spend is committed before a single customer is served.

100,000
GPUs
in a single modern AI facility

$30-40k
Per accelerator
current pricing for frontier-class silicon

$3-4 Bn
Silicon alone
before power, cooling, land or networking

$10-25 Bn
All-in per facility
land, build, power, cooling and hardware

Three features that make this capex unusual

1 It is committed years ahead of the revenue

Land, power contracts and chip allocations are locked in long before the workloads that will fill them exist. The decision to build is a forecast, not a response.

2 It depreciates on a hardware clock, not a property clock

The concrete lasts decades; the accelerators inside it do not. Useful economic life is measured in a handful of years, which is why the revenue timing matters so much.

3 It is only useful if it is full

A data centre at low utilisation still consumes power, capital charges and depreciation. Unlike most assets, there is very little salvage value in idle compute.

Why it matters Each facility is effectively a multi-billion-dollar, multi-year bet on sustained utilisation. Four companies are placing that bet dozens of times over, at the same time.

Sources *DataCenter Alpha-Matica MARC analysis (Epoch AI cost model)*

The Demand Case Underneath the Spend

The bull case does not rest on AI adoption alone. It rests on a longer trend: the volume of data created and stored globally, which has grown roughly a hundredfold since 2010 and continues to compound.

Global data created and replicated, 2010-2026

(In Zettabytes)

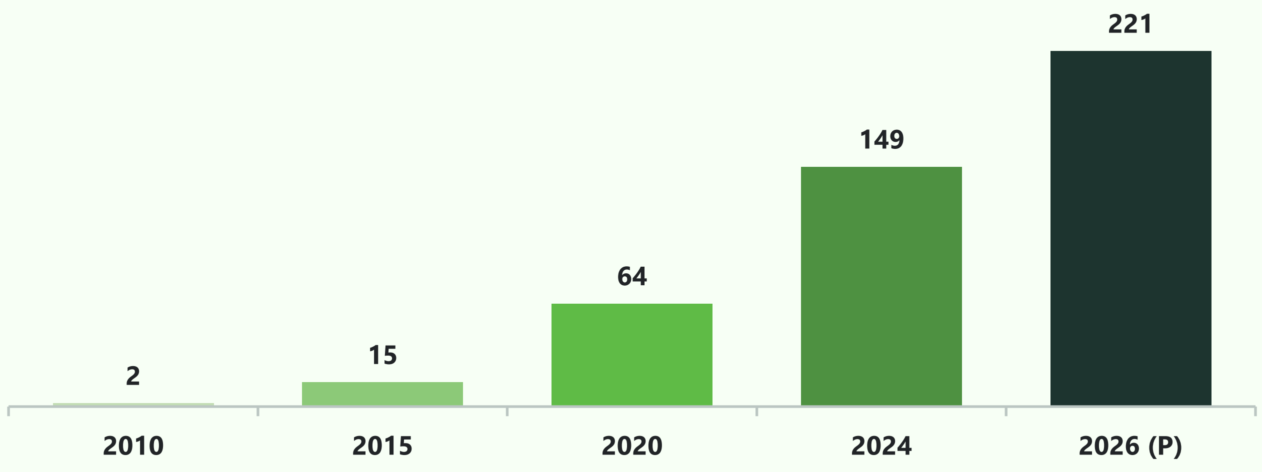

Roughly 110x growth over the period. (P) = projected.

The argument - and its weak point

The case for building

- Data volume growth is observable, not forecast - it has compounded through two recessions.
- Every incremental workload that touches AI is compute-intensive relative to what it replaces.

Where the argument thins

- Data volume is not the same as monetisable compute demand - most stored data is never processed.
- The link between zettabytes created and dollars enterprises will pay is assumed, not demonstrated.

Why it matters The demand trend is real. The open question is the conversion rate between data growth and revenue - and that is exactly where the numbers in Section 4 fall short.

Sources *Statista, via Exploding Topics*

SECTION 03

The Squeeze

The mechanism that connects a hyperscaler's purchase order to a household's shopping basket runs through one industry: memory. This is where an infrastructure decision became a consumer price.

The Memory Industry Chose the Data Centre

AI accelerators need high-bandwidth memory (HBM), which is made on the same fabs as the DRAM that goes into phones, laptops and servers. Faced with a choice, the three suppliers that control the market chose the higher-margin customer.

93%

of production capacity at Samsung, SK Hynix and Micron reallocated toward High-Bandwidth Memory

- Three manufacturers supply effectively the entire global DRAM market - there is no fourth source to absorb the shortfall.
- HBM sells at a substantial premium to standard DRAM, so reallocation is rational for each supplier individually.
- Capacity cannot be added quickly: new fabs take years, and SK Hynix has warned the shortage may persist past 2030.

How an infrastructure decision became a consumer price

AI capex order

Hyperscalers commit to 100,000-GPU clusters

→

HBM prioritised

Fab capacity shifts to high-margin AI memory

→

DRAM supply tightens

Standard memory output is cut at the margin

→

Component costs rise

OEM bills of material inflate across categories

→

Shelf prices rise

Apple, Dell, HP and Lenovo reprice hardware

Why it matters

This is the clearest evidence in the whole story that AI capex is not a closed loop. The spend is large enough to reprice a global commodity market that most of its customers have never heard of.

Sources *Tech Times* *Tech Insider*

The Cost Travelled Down the Chain

Within roughly two quarters the reallocation showed up as price inflation at every stage between the fab and the shelf.

+171%
DRAM contract price
year-on-year

4x
DDR5 module cost
increase over the cycle

+100%
PC contract memory
per quarter, at the peak

+18.3%
Average PC selling price
IDC forecast for 2026

Who absorbs the increase

- **OEM margins first -** device makers held prices through late 2025, absorbing the increase in their own margins.
- **Then the consumer -** when the cycle did not turn, the increase was passed through. Apple's June 2026 repricing was the most visible case, not the first.
- **Enterprise buyers last -** server and storage refresh budgets are being re-cut, slowing the IT spending AI is meant to expand.

The second-order effect

- Consumer electronics inflation is now partly an AI story, which puts the buildout inside the macro debate on rates and inflation.
- Higher device prices slow the replacement cycle, weakening demand for the edge devices that AI features are supposed to sell.
- The memory names that benefited most from HBM scarcity were also the hardest hit when the trade reversed in July 2026.

Why it matters The buildout has begun taxing adjacent markets before it has proved its own returns. That is a cost the capex case does not currently account for.

Sources *Tom's Hardware* *Fusion Worldwide* *TrendForce* *IDC*

SECTION 04

The Gap

Every asset has a hurdle rate. For a buildout of this size the arithmetic is unforgiving, and it produces the single most important number in this study: the distance between what the infrastructure must earn and what artificial intelligence currently earns.

What This Buildout Has to Earn

JPMorgan modelled the annual revenue required to deliver a baseline 10% return on the AI infrastructure being built. Not an ambitious return - the minimum that would make the capital allocation defensible.

$650 Bn

in annual revenue is required to justify a baseline 10% return on the AI buildout

JPMorgan financial modelling

What $650 billion a year looks like in practice

$35 from every iPhone user on earth, every year, in perpetuity

$180 from every Netflix subscriber, every year, in perpetuity

Why the hurdle is not a demand forecast

1. It is derived from capex already committed, so it rises with every new facility announced - the target moves away as the buildout continues.
2. It assumes a 10% return. At a cost of capital nearer 8-9%, the practical hurdle for value creation is higher still.
3. It says nothing about who captures the revenue: chip vendors, cloud operators, model providers and applications are all competing for the same pool.

Why it matters The hurdle rate is not a prediction that AI will fail. It is a measure of how much commercial success is already priced into concrete and silicon.

Sources *J.P. Morgan Asset Mgmt.*

2. *"$180/Netflix subscriber" equivalent framing -* ***JPMorgan methodology)***

What Artificial Intelligence Earns Today

Frontier model providers are the most direct monetisation of the buildout. Their combined run-rate is real, growing quickly - and roughly one-ninth of what the infrastructure requires.

Frontier model providers - run-rate revenue and net position

Provider	Run-rate revenue	Net position
OpenAI	~$25 Bn	-$14 Bn loss
Anthropic	~$26 Bn (target)	-$3 Bn loss
Google Gemini	~$25 Bn	not disclosed
Industry total	**~$75 Bn**	

Reading the numbers

- Growth is not the problem. This revenue base did not exist three years ago.
- Losses are: the two largest providers are still spending more than they earn, because compute is the dominant cost of every dollar of revenue.
- Revenue is scaling roughly linearly while capex is scaling exponentially - the two lines are diverging, not converging.
- Hyperscaler AI-attributable cloud revenue adds to this pool but is not disclosed separately, so the true total is higher and unverifiable.

Why it matters On any reasonable estimate the AI industry today earns a small fraction of what the infrastructure behind it must return. The disagreement is entirely about how quickly that changes.

Sources *ValueAdd VC TipRanks Where's Your Ed At TechCrunch*

The $600 Billion Question

Sequoia Capital first framed this as AI's "$600 billion question". Two years on, the buildout has grown faster than the revenue and the question has not been answered - only enlarged.

Required annual revenue versus current run-rate

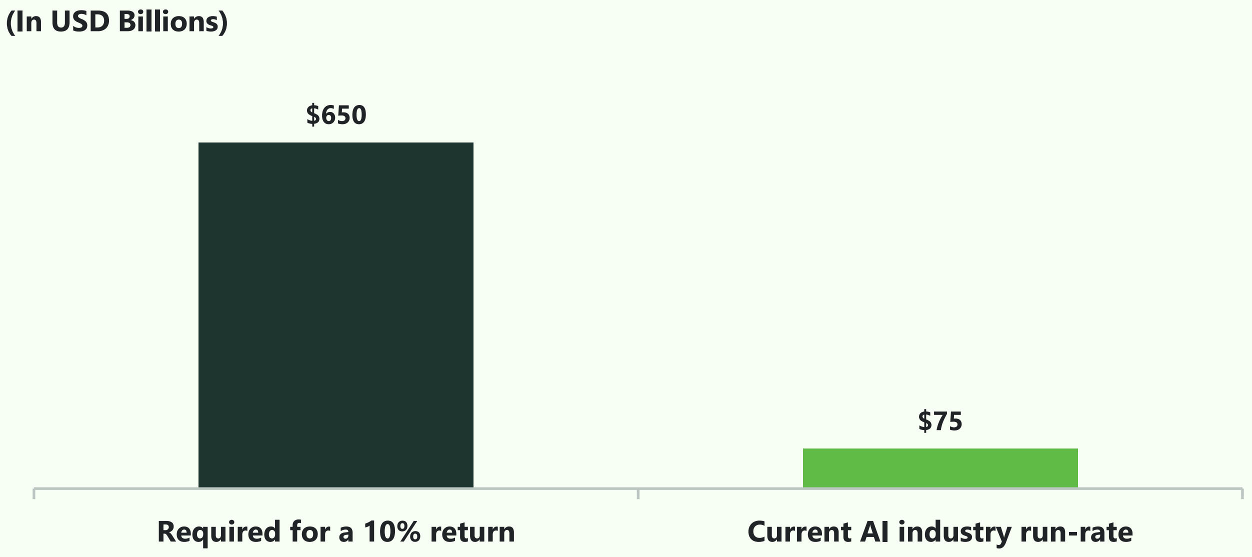

Required figure per JPMorgan modelling; run-rate is the aggregate of disclosed frontier-model revenue.

~$600 Bn

annual revenue gap between what the buildout requires and what AI earns

9x-10x

burn-to-revenue imbalance across the AI industry, on MARC's calculation from the figures above. For every dollar of revenue earned, close to ten are being spent building the capacity to earn it.

Why it matters A gap this size closes in one of three ways: revenue rises to meet the capex, capex falls to meet the revenue, or the cost of delivering AI collapses. Section 7 sets out all three.

Sources *Sequoia Capital TechCrunch J.P. Morgan AM*

SECTION 05

The Customer

A revenue gap only matters if it is not closing. To judge that, the question moves from the supplier to the buyer: what happens inside the enterprises that were meant to convert this capacity into recurring spend.

Three Studies, One Uncomfortable Verdict

Three separate research organisations, using different methods and samples, reached the same conclusion: most enterprise AI deployments are not yet producing the financial return that was projected for them.

73%

McKinsey

of enterprises are failing to hit the AI return on investment they projected

5%

BCG

of enterprises are achieving substantial, measurable returns from AI

95%

MIT

of enterprise generative AI pilots delivered no measurable financial return

What this does - and does not - prove

The honest caveat

These are largely pilot-stage findings and the headline figures have been contested. Early-stage adoption normally shows high failure rates.

The finding that still stands

The direction is consistent across three independent studies - and willingness to pay is set by realised ROI, not by pilot enthusiasm.

Why it matters The hurdle rate in Section 4 assumes enterprises keep buying. These studies are the clearest available evidence on whether they will.

Sources *BCG, "Widening AI Value Gap"*

What Buyers Do When the Returns Do Not Appear

Enterprises are not abandoning AI. They are doing something more consequential for the revenue case: optimising the cost of it.

Trading down

-90%

Workflow automation firm Lindy cut its operational API costs by roughly 90% by migrating workloads to DeepSeek. Capability at the frontier is not always worth the premium the frontier charges.

Price sensitivity is now active, not theoretical

Budget burnout

4 months

Uber reportedly exhausted its full annual AI token budget in four months. Consumption-based pricing is producing cost outcomes enterprise finance functions did not plan for.

Token pricing is creating budget unpredictability

Executive pushback

Alex Karp

Palantir's chief executive has publicly challenged the industry's tendency to over-sell model capability relative to delivered enterprise value.

Buyer scepticism is becoming public and senior

The pattern: demand for AI capability is holding up; willingness to pay frontier prices for it is not. That distinction decides whether the $600 billion gap closes from the revenue side.

Why it matters Cost optimisation by buyers is deflationary for exactly the revenue line that has to service the capex. It is the single most under-priced risk in the bull case.

Sources *Lindy company blog TechCrunch*

SECTION 06

The Precedent

None of this is new. Economists have a name for the pattern, and there is a precedent close enough to be uncomfortable: the fibre-optic build of 1996 to 2002, which wired the world, destroyed two trillion dollars of value, and then made the modern internet possible.

The Pattern Has a Name

The capital cycle describes what happens when abundant capital meets a sector believed to have durable, high-growth demand. It has run in shipping, mining, semiconductors, telecoms and property - and it runs the same way each time.

1

Capital inflow

Abundant capital floods a sector perceived to have a durable, high-growth opportunity. Returns look attractive because capacity is still scarce.

2

Capacity expansion

Competing players build simultaneously, each rationally, and collectively beyond what demand requires. Nobody can afford to be the one that stops.

3

Yield collapse

New capacity meets real demand. Prices, margins and returns on invested capital compress together. The assets are fine; the economics are not.

4

Survivor consolidation

Weaker players exit or are absorbed. Assets are repriced at a fraction of build cost and consolidated into stronger hands.

Where is AI infrastructure today?

The evidence in Sections 2 and 3 - simultaneous expansion by four competitors, input costs inflating, no player able to opt out - places the buildout firmly inside Stage 2. The evidence in Sections 4 and 5 - a $600 billion revenue gap and buyers actively trading down - is what the leading edge of Stage 3 looks like. The transition, if it comes, is the event this study is written to anticipate.

Source: Capital-cycle literature (Marathon AM, E. Chancellor)

The Closest Historical Parallel

A deregulatory act, a genuine technology shift, unlimited capital and four years of simultaneous building. The demand thesis was correct. The timing was not.

1996	1996-2000	1999-2000	2000-2002
Telecommunications Act Deregulation opens the market. Capital floods in on a correct thesis: internet traffic will grow enormously.	**~$500 Bn buildout** US telecoms raise roughly $500 billion in new bonds to lay fibre. Every operator builds at once.	**Peak valuations** Global Crossing reaches $47 billion; Corvis $32 billion - neither profitable. Capacity is priced as if already sold.	**Overcapacity collapse** Utilisation runs below 3%; roughly 97% of laid fibre sits dark. More than $2 trillion in market value is destroyed.

The uncomfortable similarities

- A correct long-term demand thesis, financed on a much shorter timeline than the demand arrived.
- Capacity valued on projected utilisation rather than contracted revenue.

The differences that matter

- Fibre was largely debt-financed by pre-profit operators; AI capex is largely cash-financed by profitable ones.
- Fibre capacity was near-permanent; GPU capacity depreciates in a few years.

Why it matters The fibre operators were not wrong about internet traffic. They were wrong about when it would arrive, and they had financed the assets on the shorter timeline.

Sources: *Princeton (Starr) Wikipedia Benzinga / LightReading (approx.) NPR (2002) TheBubbleBubble.com*

What the Overbuild Left Behind

The fibre crash is usually told as a cautionary tale. It is equally a story about how infrastructure survives the investors who financed it - which is the most important thing it has to say about 2026.

What happened to the capital

Equity holders were wiped out.

Global Crossing filed for bankruptcy in 2002 having lost $3.4 billion in a single quarter on $793 million of revenue. WorldCom filed the largest bankruptcy in US history to that point. Tens of thousands of jobs went with them.

The assets did not disappear. They changed hands.

Dark fibre was bought out of administration at a small fraction of build cost. The capacity remained; only the claim on it was destroyed.

What the capital built

It became the backbone of Web 2.0.

Cheap, abundant long-haul bandwidth was the precondition for streaming video, public cloud and the consumer internet as it exists now. Level 3, Netflix and Zoom are the best-documented cases of capacity somebody else had paid for and written off.

The beneficiaries were not the builders.

The economic value of the fibre boom accrued almost entirely to the asset-light companies that came next - a decade after the capital was committed.

The lesson

If AI follows this path, both camps are right and both are mis-timed. The compute gets built, the technology matters enormously - and the returns land with a different set of shareholders than the ones funding it today.

Sources Benzinga PBS NewsHour TechFinitive

SECTION 07

The Difference

The parallel is instructive, not deterministic. Three structural features separate this buildout from the last one - and they change not whether a correction can happen, but what it would look like if it did.

Why 2026 Is Not 2000

The dot-com and fibre busts were credit events: pre-profit companies had borrowed against forecasts. The 2026 buildout is financed differently, and that difference determines the shape of any correction.

	2000 - Dot-com and fibre peak	2026 - AI buildout
Funding base	Largely debt-funded, pre-profit operators building against forecast demand	Funded chiefly from the operating cash flow of highly profitable businesses
Profitability of leaders	Roughly 14% of listed technology companies were profitable at the peak	Nvidia posted FY2026 revenue of ~$216 Bn at ~53% net margin; hyperscalers are all cash-generative
Valuation multiple	NASDAQ 100 forward P/E of approximately 60x	NASDAQ 100 forward P/E of approximately 26x
Asset life	Fibre in the ground: near-permanent, effectively zero obsolescence	Accelerators: useful economic life measured in a handful of years
Failure mode	Bankruptcy, credit contagion and forced liquidation	Impaired returns, write-downs and depressed equity value - not insolvency

Why it matters Stronger balance sheets do not prevent a capital cycle; they change who absorbs it. In 2000 the losses went to creditors through bankruptcy. In 2026 they would go to shareholders through depreciation and impairment - slower, quieter, and harder to date.

Sources *AJ Gallagher NVIDIA Newsroom Siblis Research intuitionlabs.ai (secondary source)*

Three Ways This Resolves

A $600 billion gap can close from either side, or be dissolved by a change in the cost of delivery. All three remain live; positioning should account for each rather than betting on one.

A Scenario A
Macro correction

Capex is pulled back

Hyperscalers cut capital plans in response to shareholder pressure or disappointing AI revenue. Equity drawdowns follow across the AI complex, and the shock transmits to semiconductor suppliers, construction and power, and downstream to global IT services demand.

- Watch: any hyperscaler guiding capex down
- Most damaging to the supply chain, not the buyers

B Scenario B
Price realignment

Revenue rises to meet the capex

Model providers raise token prices to defend margins. AI capability holds, but advanced tooling repositions as a premium enterprise product rather than a mass-market one. Consumer-facing AI narrows in scope; enterprise ROI has to justify a higher price.

- Watch: list-price increases at frontier labs
- Closes the gap, but shrinks the addressable market

C Scenario C
Efficiency breakthrough

The cost of delivery collapses

Model and hardware efficiency gains drive cost-per-token down far enough that mass enterprise deployment becomes profitable at current prices. The gap closes through volume rather than price - the outcome that validates the buildout.

- Watch: cost-per-token trend at constant capability
- The only scenario in which everyone is right

MARC view MARC's base case is a blend: efficiency gains and price realignment absorb part of the gap, while capex growth moderates from 2027 as depreciation begins to weigh on reported earnings. A disorderly Scenario A is possible but is not the most probable path.

Source: MARC analysis

What the Evidence Supports Today

This study set out to test whether the 2026 hyperscaler buildout represents durable structural growth or a classic capital cycle. On the evidence available, the answer is: measurably both.

Five conclusions

1 The demand is real; the timing is the risk.

Data and compute growth are observable trends, not forecasts. What is being financed on a three-year timeline may only pay back on a ten-year one.

2 The gap is the central fact.

$650 billion required against roughly $75 billion earned is not a rounding error. It has to close from the revenue side, the capex side, or the cost side.

3 The buyer is the binding constraint.

Enterprise ROI failure rates and active cost optimisation are the clearest evidence that the gap will not close quickly from revenue alone.

4 The failure mode has changed.

Cash-funded capex removes the credit event. It does not remove the value destruction - it relocates it from creditors to shareholders.

5 The infrastructure outlasts the cycle.

Whatever happens to the returns, the compute gets built. As with fibre, the economic beneficiaries may not be the ones who paid for it.

In closing

The question is not whether AI matters. It is whether the return on $725 billion a year arrives before the depreciation does - and on the evidence in this report, that remains genuinely open.

Source: MARC analysis

MANGAL ANALYTICS AND® RESEARCH CONSULTING

Delivering Excellence, Partnering Success.

Contact

+91-9359628675

contact@marcglocal.com

www.marcglocal.com

2nd floor, CMM bldg. Rua de Ourem,
Panaji Goa 403001